GO-TO-MARKET DOCUMENT

# Evenline: the AI receptionist that learns your business from your website

A phone receptionist for small businesses that answers every call, takes the details that matter, and texts the owner a summary minutes later — set up in about a minute from a website address.

| Document | Go-to-market plan and cost management |
|---|---|
| **Product status** | Working build; two onboarding paths; 1,043 backend and 177 frontend tests passing |
| **Deployment** | AWS us-east-2 · ECS Fargate, ALB, CloudFront, RDS, ElastiCache |
| **Prepared** | 14 October 2026 |

## Executive summary

Small businesses lose revenue to unanswered calls, and the staff who would answer them are the same staff serving customers in front of them. Evenline answers every call with a voice agent that already knows the business, captures who called and why, and reports back by email and dashboard. It is live on AWS today.

The wedge is onboarding. Competing products ask an owner to fill in a long form before they hear anything work. Evenline reads their existing website, shows them what it learned, and lets them correct it — the first interaction is a review, not a questionnaire.

> **The two numbers this document exists to answer**
>
> **Deployment cost** — roughly **$3.00 to $4.60 per day** ($90–140 per month) to run the whole platform on AWS at current scale, dominated by three fixed line items. **Tooling and API cost** — roughly **$1.20 to $3.60 per day** during the build, almost all of it vendor usage rather than software licences. Both are broken down in sections 6 and 7, with the assumptions stated.

**What is ready:** both setup paths, provisioning, live call answering, transcripts and summaries, the customer dashboard, deletion with number release, and calendar integrations. **What is not:** published prices, outbound SMS sending, and languages other than English. Each is called out where it affects the plan rather than left implied.

# 1. The product

What a customer gets, and what happens on a call.

## Two ways to set up

Both paths end at the same provisioning pipeline, so there is one way for a receptionist to be created and one place where money is spent.

| Path | What the owner does | Time |
|---|---|---|
| **Use my website** | Enters a website address. The platform reads the public pages, drafts what the receptionist should know, and shows it for review and editing before anything is created. | ~1 minute |
| **Set up manually** | Fills in business details, services, hours, greeting and escalation rules across four steps. | ~5 minutes |

## On a call

A caller dials the business number. The voice agent answers with the greeting the owner chose, answers questions from the knowledge the owner approved, collects the details the owner asked for, and escalates according to the owner's rules. After the call, a transcript and a short summary with an urgency rating arrive by email and appear on the dashboard.

## In the dashboard

- **Home** — the number, usage against plan, and the latest calls.
- **Calls** — every call with transcript, summary, intent and callback number.
- **Agent** — what the receptionist knows, edited and republished as a new version; every past call stays pinned to the version that answered it.
- **Contacts** — callers, grouped, with their history.
- **Integrations** — Google Calendar and Microsoft Outlook via real OAuth, filtered by business type.
- **Settings** — business details, notifications, phone, billing, and deleting the receptionist.

## Verticals supported today

| Vertical | Why it fits |
|---|---|
| **Salon & spa** | High call volume during appointments; staff hands are literally full. |
| **Law firm** | Intake questions are structured and repetitive; missed calls are expensive. |
| **Medical practice** | Front desk saturation; integrations deliberately exclude general automation tools for privacy reasons. |
| **Real estate** | Agents are mobile and miss inbound enquiries. |
| **Something else** | Generic configuration, same pipeline. |

**Honesty note carried into the product:** the dashboard shows “Not tracked yet” for hang-ups, spam and transfers rather than displaying a zero, because those are not measured yet. That principle should carry into sales collateral too.

# 2. Market and ideal customer

Who feels this pain enough to pay monthly for it.

## Ideal customer profile

| Attribute | Target |
|---|---|
| **Business size** | 1–20 staff, single or few locations |
| **Call volume** | 20–400 inbound calls a month |
| **Current solution** | Voicemail, a shared mobile, or an answering service at $200–600/month |
| **Buying trigger** | Missed calls during busy periods; a receptionist leaving; opening a second location |
| **Decision maker** | Owner or practice manager — one person, no procurement |
| **Has a website** | Yes — this is the qualifying signal for the fast onboarding path |

## Why they buy

- **A missed call is a lost job.** For a salon or a law firm, one booking covers a month of subscription.
- **Answering services cost more and know less.** A human service reads a script; this reads the business's own website.
- **Voicemail is worse than nothing.** Most callers hang up and call the next listing.

## Competitive position

| Alternative | Their edge | Our counter |
|---|---|---|
| **Human answering service** | Real people, nuance | A fraction of the price, answers instantly, never busy, and knows the actual services and prices |
| **Upfirst and similar AI receptionists** | Funded, further ahead on polish | Equivalent website onboarding, plus honest reporting and per-business knowledge the owner edits directly |
| **Voicemail / nothing** | Free | Quantified: every missed call is visible in the dashboard the day they start |
| **Generic AI agent builders** | Flexible | Vertical defaults and real telephony provisioning in one step, not a toolkit |

## Positioning statement

**One sentence**

For small businesses that lose work to unanswered calls, Evenline is an AI receptionist that learns the business from its own website in about a minute — so the first thing an owner sees is their own services read back to them, not a blank form.

# 3. Packaging and pricing

The plans exist in the product. The prices do not, and that is the next decision.

## Plans as built

| Plan | Included minutes | Intended buyer | Price |
|---|---|---|---|
| **Starter** | 500 / month | A single location that wants every call answered | Not set |
| **Pro** | 2,000 / month | Busy teams with a steady stream of calls | Not set |
| **Enterprise** | 10,000 / month | Multi-location, high volume | Not set |

Prices are deliberately absent from the codebase so that no figure is shown to a customer before it is decided. Setting them is a one-line change per plan.

## What the floor has to cover

Pricing must clear variable cost per minute plus a share of fixed infrastructure. Using the per-call economics in section 5, a Starter customer consuming half their included minutes costs single-digit dollars a month in vendor spend. The constraint is not cost — it is positioning against the $200–600 answering service the customer is comparing against.

## Recommended structure

- **Anchor against the alternative, not against cost.** The comparison is a human answering service, not an API bill.
- **Charge monthly with included minutes and a clear overage rate** — the usage meter and plan limits are already built and enforced.
- **Keep a 14-day trial**, which the billing gate already grants, but see the warning below.
- **Do not discount the first month.** The value is proven in week one by the Calls page; a discount invites churn at renewal.

> **Commercial risk to close before launch**
>
> The deployed environment currently grants a trial on signup with live vendors enabled, which means **every signup at the public URL buys a real phone number** that bills monthly until released. Before any public launch, either require a payment method at signup or cap trial provisioning. This is the single largest uncontrolled cost in the product today.

## Expansion revenue, in build order

| Lever | State | Why it earns its place |
|---|---|---|
| **Calendar booking during a call** | Integrations connected; booking not built | The most requested capability; turns a message-taker into a booker |
| **SMS follow-up** | Templates configurable, sending not built | Caller asks for a link; agent texts it |
| **Multi-location / multi-number** | Data model supports it | Natural Enterprise upsell |
| **Languages beyond English** | Not supported | Opens whole regions; requires provider work |

# 4. Go-to-market motion

How the first hundred customers arrive.

## Channel plan

| Channel | Motion | Why it suits this product |
|---|---|---|
| **Founder-led outbound** | Call the business. Let their own receptionist answer. Leave the demo as the voicemail. | The product demonstrates itself in thirty seconds, in their own words |
| **Vertical partnerships** | Salon software, practice management, local bar associations | One integration or referral reaches hundreds of qualified owners |
| **Local search and directories** | Rank for “answering service” intent by city and vertical | Buying intent is explicit and geographic |
| **Referral** | Owners talk to other owners in the same trade | Trust transfers inside a vertical better than any ad |

## The demo that sells it

Take the prospect's website address during the conversation. Paste it in. Thirty seconds later, read their own services and hours back to them from the review screen. The demo is the onboarding, which means every demo is already half an installation.

## Launch sequence

| Phase | Goal | Gate to the next phase |
|---|---|---|
| **0 — Close the cost risk** | Payment method or trial cap before public signup | No signup can provision a number unattended |
| **1 — Design partners (5–10)** | One vertical, hand-held onboarding, daily feedback | Call containment and summary quality judged acceptable by owners |
| **2 — Paid pilot (25–50)** | Prices set, trials converting | Conversion and churn measurable over two cycles |
| **3 — Scale one vertical** | Repeatable acquisition in a single trade | CAC payback under three months |

## Metrics that decide whether this works

- **Setup completion rate**, website path versus manual — the core bet of the product.
- **Time from signup to first answered call.**
- **Calls handled without escalation**, by vertical.
- **Summary accuracy**, sampled and read by a human weekly — the thing that destroys trust fastest.
- **Trial-to-paid conversion**, and **gross margin per customer** once prices exist.

# 5. Cost management I — deployment

What it costs to keep the platform running on AWS, per day.

Built from the deployed task definitions and architecture in `infra/ecs/`. Live billing could not be read while preparing this document because the AWS session had expired, so these are list-price estimates, not invoiced amounts. The verification commands are on page 8.

## Running cost by component

| Component | Specification | Per month | Per day |
|---|---|---|---|
| **ECS Fargate — API** | 0.5 vCPU, 1 GB | $18 | $0.59 |
| **ECS Fargate — frontend** | 0.25 vCPU, 0.5 GB | $9 | $0.30 |
| **ECS Fargate — worker** | 0.25 vCPU, 0.5 GB | $9 | $0.30 |
| **ECS Fargate — Temporal worker** | desired count 0 | $0 | $0.00 |
| **Application Load Balancer** | hourly + ~1 LCU | $18–22 | $0.60–0.72 |
| **RDS PostgreSQL** | single-AZ, class to confirm | $15–28 | $0.49–0.92 |
| **ElastiCache Serverless** | minimum spend | $6–60 | $0.20–1.97 |
| **Public IPv4 × 3 tasks** | $0.005 per hour each | $11 | $0.36 |
| **CloudWatch Logs, ECR, Secrets Manager** | 1 secret, few images | $4 | $0.13 |
| **CloudFront** | low traffic, free tier | $0–1 | $0.00–0.03 |
| **Total** | **one task per service** | **$90–140** | **$3.00–4.60** |

## Where the money actually goes

- **ElastiCache Serverless is the one to check first.** It bills a minimum whether used or not, and Redis here is only a cache — every read already falls back to PostgreSQL. A small node, or disabling it, is the largest single saving available.
- **The load balancer is pure fixed cost** at roughly $20 a month regardless of traffic. Unavoidable while HTTPS is served through CloudFront.
- **Public IPv4 addresses cost about $11 a month** because tasks run in public subnets. This was the right trade: a NAT Gateway would cost more.
- **Temporal is costing nothing** — the orchestrator runs as a state machine on the existing worker, with the Temporal service scaled to zero.

> **Scaling shape**
>
> Of the $3.00–$4.60 a day, roughly **$2.70 is fixed** — it does not change whether the platform serves one customer or two hundred. Per-customer marginal infrastructure cost is close to zero until call volume forces a second API task. That makes early customers almost entirely gross-margin positive once vendor usage is covered.

# 6. Cost management II — tools and APIs

What the third-party services cost while building and running the product, per day.

## Build-time cost, per day

These are the services consumed during development: model calls to generate agent configurations and website knowledge, telephony numbers bought for testing, and voice minutes spent on test calls. Rates are list prices and should be confirmed against each vendor's invoice.

| Service | What it is used for | Basis | Per day (build) |
|---|---|---|---|
| **OpenAI (gpt-5.4-mini)** | Agent configuration, call summaries, website knowledge extraction | Per token; tens to low hundreds of calls a day in development | $0.30–1.50 |
| **Twilio** | Phone numbers and inbound voice minutes | ~$1.15 per number per month, plus per-minute voice | $0.10–0.50 |
| **ElevenLabs** | The voice agent itself — the dominant per-minute cost | Per conversation minute | $0.50–1.50 |
| **Resend** | Call summary and system email | Free tier covers development volume | $0.00 |
| **Temporal Cloud** | Not in use; state-machine orchestrator instead | — | $0.00 |
| **GitHub, Docker Desktop, ngrok** | Source control, builds, local webhook tunnel | Free tiers at this team size | $0.00 |
| **Total** | **build-time tooling and APIs** | | **$1.20–3.60** |

**Important:** most development runs in dry-run mode, where every vendor is replaced by a fake and the true cost is zero. The figures above describe days when live vendors were enabled for end-to-end testing. The entire automated test suite — 1,043 backend tests — runs against fakes and costs nothing.

## Per-call unit economics, in production

This is the number that decides gross margin. A typical answered call runs two to four minutes.

| Cost element | Driver | Per 3-minute call |
|---|---|---|
| **Voice agent minutes** | ElevenLabs conversation time | $0.24–0.36 |
| **Telephony minutes** | Twilio inbound voice | $0.03 |
| **Summarization** | One model call per finished call | $0.01 or less |
| **Number rental, amortized** | One number per customer | $0.04 per day regardless of calls |
| **Marginal cost per call** | | **~$0.28–0.40** |

A Starter customer using 250 of their 500 included minutes costs roughly **$25–$35 a month** in vendor spend. Any monthly price set meaningfully below the $200–600 answering service it replaces still leaves a healthy margin — provided the trial cannot provision numbers unattended.

# 7. Cost controls, risks and next steps

What keeps the bill honest, and what to do first.

## Controls already built into the product

- **Dry-run mode** replaces every paid vendor with a fake, so development and the whole test suite cost nothing. Crawling is bounded too — page count, depth, bytes, time and concurrency are capped, so no website can run up an open-ended model bill.
- **A billing gate before the money step.** Provisioning pauses before a number is bought unless the account is entitled, then resumes when billing clears.
- **Idempotent purchasing.** Three guards — an existing active row, a vendor lookup by name, and a pending row written before the call — make buying two numbers for one business structurally hard. A setup that fails terminally releases its number and deletes its agent rather than leaving both billing silently.
- **Deletion with release.** Owners delete a receptionist and hand the number back in one action, stopping the rental.

## Open cost risks

| Risk | Exposure | Action |
|---|---|---|
| **Trial signups buying real numbers** | $1.15/month each, indefinitely | Require a payment method, or cap trial provisioning, before launch |
| **Orphaned numbers from testing** | Unknown until audited | Audit Twilio for numbers with no active business; release them |
| **ElastiCache Serverless minimum** | Up to ~$60/month for a cache that has a database fallback | Measure usage; move to a small node or disable |
| **No spend cap on model usage** | Grows with call volume | Record token usage per call — the ledger exists, unused — and alert on anomalies |

## Verifying these numbers

Every figure in sections 5 and 6 is a list-price estimate. To replace them with invoiced amounts, authenticate and run:

| Question | Command |
|---|---|
| **What did AWS actually cost?** | `aws ce get-cost-and-usage --granularity MONTHLY --metrics UnblendedCost --group-by Type=DIMENSION,Key=SERVICE` |
| **How many tasks are running?** | `aws ecs describe-services --cluster ai-receptionist --services ...` |
| **What are the database and cache?** | `aws rds describe-db-instances` and `aws elasticache describe-serverless-caches` |

## Next three actions

- **Close the trial provisioning gap.** It is the only uncontrolled spend in the product and it blocks any public launch.
- **Set prices.** Everything else in the commercial model waits on this; the plans already carry them.
- **Run the billing audit above** and replace these estimates with invoiced figures before quoting margins outside the team.

Prepared from the repository and deployment as of 14 October 2026. Product facts come from the codebase; monetary figures are list-price estimates.